

# Exercising Right: Using sensors to ensure proper form

michael downs

## synopsis

Many people measure the amount they exercise (miles run, weights lifted, etc.). Few measure the quality of their exercise, their “form”. Using data from sensors placed on the belts, forearms, arms, and dumbbells of six participants, I used four algorithms to classify barbell lifts. Classifications range from correct form (class A) to four incorrect forms including throwing the elbows (class B), throwing the hips (class E), etc..

I found random forest (rf) and boosting (gbm) to be effective at classifying activities based on sensor output. Both algorithms achieved 97%+ sensitivity and 99%+ specificity scores on cross validation. (The random forest model scored 20/20 on the project test set.) Further, both models maintained high cross validation scores: 1. across four train/test datasets that were constructed using different methodologies and 2. across a range of variables from over 100 down to as small as 10. K-nearest neighbors (knn) and basic regression trees (rpart) were less successful with sensitivity scores of  $\approx 88\%$  and  $\approx 53\%$  on cross validation, respectively.


I walk thru model construction below including use of cross validation, estimates of out-of-sample error rates and the rationale for decisions made along the way.

## the model

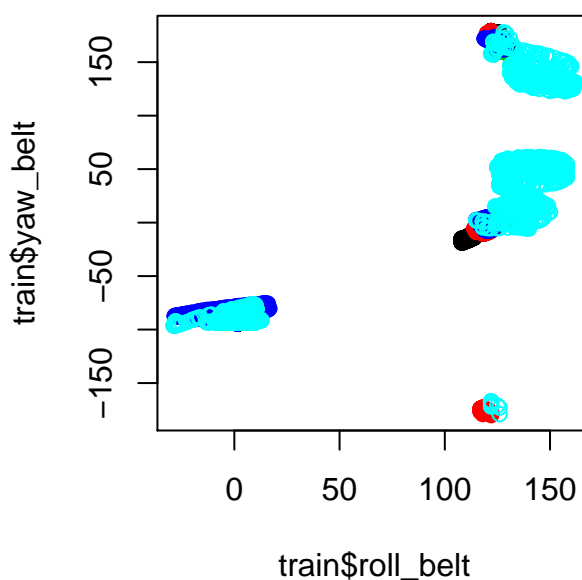
### getting data

The training data set included 160 variables and about 20,000 records. Each record contained a subset of sensor readings sampled during the exercise period for each of the six participants. Basic exploratory analysis showed that 67% of the columns contained 97% NA values which would have to be cleaned. That analysis also identified a roughly even distribution of observations among the three classes and significant signal that could be seen when fields were plotted against each other (below).

**class distribution**



**clearly signal here**

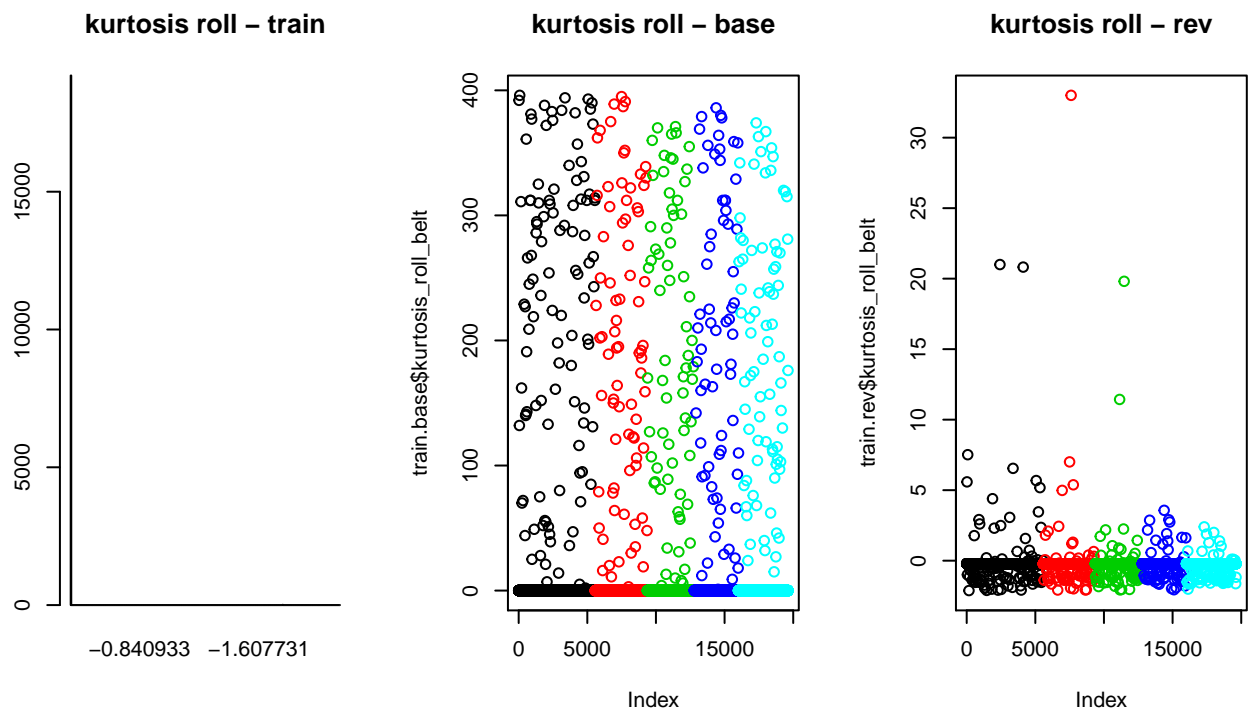


## cleaning data

I generated four train/test datasets. While I deleted timestamp and window columns for all datasets (to the extent exercisers are following a script, these columns provide information that wouldn't be available otherwise), each dataset was developed according to its own methodology. Highlights include: 1. *\$train.base*: I replaced NA's by the means for their respective columns. Given most algorithms use column variance for predictions, this approach eliminates NAs while not changing column variance. 2. *\$train.nzv*: I deleted rows with near zero variance (`nzv()`) then set NAs in remaining columns to the respective column mean. 3. *\$train.na*: I deleted columns with greater than 97% NA. 4. *\$train.rev*: I converted factor to numeric using `levels(train.rev[,i])[train.rev[,i]]`, then set any remaining NAs to the respective column means.

The effect of these changes and others not covered is illustrated in the kurtosis graphic below. From left to right I show how the data was imported, what the data looks like using for *train.base* and what the data looked like for *train.rev*.

```
## Loading required package: lattice
## Loading required package: ggplot2
```



## pca pre-processing

I performed feature selection in two steps. The first was pca analysis. I based my approach on the lecture where pca was used only for highly correlated variables (rather than to reduce dimensions and regularize all variables). Specifically, I looped over the data sets five times replacing highly correlated variables with their principal components. As the approach was a bit involved, I've included the code below. This approach reduced variables between 25% and 30% across data sets.

```
dset=list(ds=list("train.base","train.nzv","train.na","train.rev"),
          trnx=list(train.base,train.nzv,train.na,train.rev),
          trny=train$classe)
```

```

par(mfrow=c(1,1))
plot(1:5,type="n",main="pca dimension reductions",ylim=c(0,160),xlab="pca iterations",ylab="total variat

for(k in 1:length(dset$trnx)){
  target=dset$trnx[[k]][,c(2:(dim(dset$trnx[[k]])[2]-1))]
  gph.line=NULL
  cat("***** new file: ",dset$ds[[k]],"\n")
  for(i in 1:5){ # was 7
    cat("1. start: dim(target)",dim(target),"\n")
    gph.line[i]=dim(target)[2]
    cor.decs=matrix(NA,dim(target)[2],2)
    M=abs(cor(target)) # find predictors w/ high correlation
    diag(M)=0 # set cor for vars w/ themselves to 0

    for(j in 1:dim(target)[2]){
      cor.decs[j,1]=j
      cor.decs[j,2]=length(which(M[,j]>max(0.5,(1-0.1*i)),arr.ind=T))
    }
    clean=target[,c(which(M[,which.max(cor.decs[,2])>max(0.5,(1-0.1*i)),arr.ind=T),
      which.max(cor.decs[,2]))]

    cat("2. consolidating: ",cor.decs[which.max(cor.decs[,2]),][2]+1,"\n")
    comps=prcomp(clean);print(summary(comps))

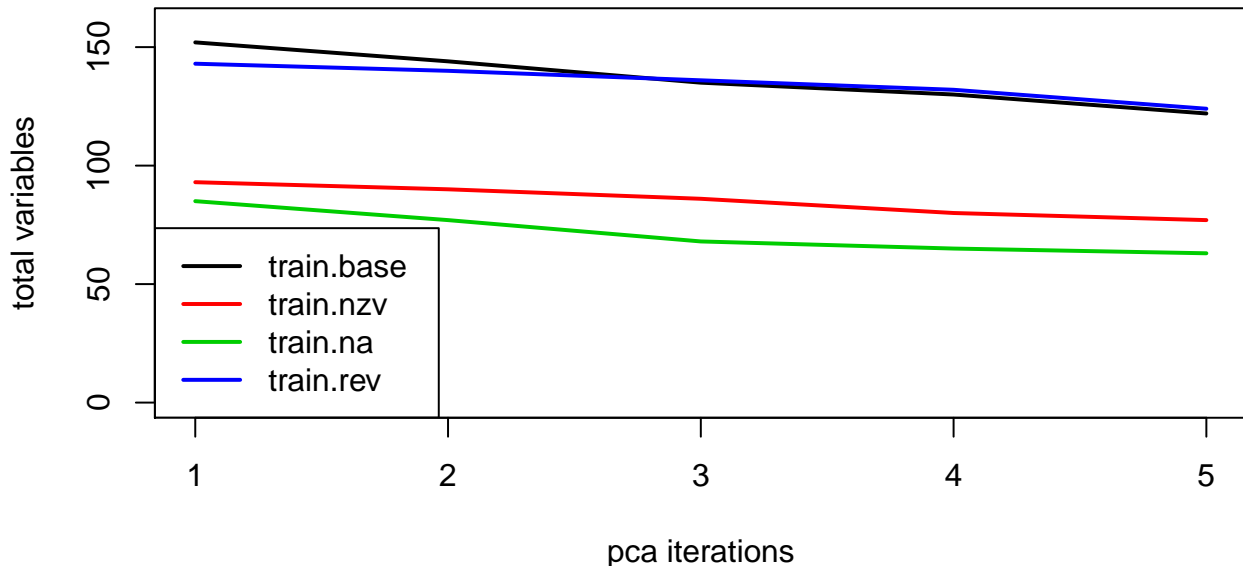
    target=target[,-c(which(M[,which.max(cor.decs[,2])>max(0.5,(1-0.1*i)),arr.ind=T),
      which.max(cor.decs[,2]))]

    tmp=as.matrix(summary(comps)$importance);pc.cut=which(tmp[3,]>0.95)[1]
    comps=comps$x[,1:pc.cut];comps=as.data.frame(comps)
    for(z in 1:pc.cut){
      colnames(comps)[z]=c(sprintf("%s_%s_%s_%s","iter",i,"pc",z))}
    cat("3. into: ",dim(comps)[2],"\n")

    target=cbind(target,comps)
  }
  dset$trnx[[k]]=target
  lines(1:5,gph.line,col=k,lwd=2)
}
legend("bottomleft",c("train.base ","train.nzv ","train.na ","train.rev "),col=c(1,2,3,4),lwd=2)

```

## pca dimension reductions



```
rm(tmp, clean, target)
for(i in 1:4){dset$trnx[[i]]$user_name=train$user_name
              dset$trnx[[i]]$classe=train$classe}
```

## models

I built a two loop model training structure to iterate thru the datasets outlined above and a set of core algorithms including regression trees (rpart), k-nearest neighbors (knn), random forest (rf) and boosting (gbm). Each iteration of the internal loop fit the model, made predictions on a cross validation set and stored off the fit for subsequent processing.

Getting thru the dataset required careful resource management. While data cleaning / pca eliminated  $\frac{1}{3}$  to  $\frac{1}{2}$  of the columns, early iterations from 70 to 115 columns. Accordingly, I trimmed train/test database by up to 50%. I also enabled parallel processing.

Beyond cross validation performed within the Caret function which showed 95-99% accuracy (1%-5% OOB error rates), I separately **cross validated** each dataset / algorithm combination using a 70 / 30 train / holdout test set cross validation.

As the approach was fairly involved, I've provided code below. Below that are the cross validated prediction sensitivity and specificity for each model. The left graphic shows all the models including the outperformance of the rf and gbm models. The right looks more closely at the rf and gbm results.

```
# controller for algo iterations
algos=list(algo=list("rpart","rf","knn","gbm")) ## logistic, lda failed.

# structures to hold rf fits
fset=list(data=list("train.base","train.nzv","train.na","train.rev"),
          rpart=list(NA,NA,NA,NA),
          rf=list(NA,NA,NA,NA),
          knn=list(NA,NA,NA,NA),
          gbm=list(NA,NA,NA,NA))
```

```

# process data sets, algos, folds
# rm(res,res.mstr)

## enable parallel processing
library(parallel); library(doParallel)
registerDoParallel(clust <- makeForkCluster(detectCores()))

for(i in 1:length(dset$ds)){
  # cut data in half
  train.cv=dset$trnx[[i]][train$user_name==unique(train$user_name)[1] |
                                                                    train$user_name==unique(train$user_name)[2] |
                                                                    train$user_name==unique(train$user_name)[3],]

  for(k in 1:length(algos$algo)){
    # train and test sets
    set.seed(123)
    inTrain=createDataPartition(y=train.cv$classe,p=0.7,list=FALSE)
    training=train.cv[inTrain,]
    testing=train.cv[-inTrain,]

    cat("***** FITTING: algo:",algos$algo[[k]],", data dim: ",dim(training),"\\n")
    # fit and predict
    fit=train(classe~,training,method=algos$algo[[k]])


    # store and print fit
    fset[[k+1]][[i]]=fit;print(fit)

    # predict
    preds=predict(fit,testing)

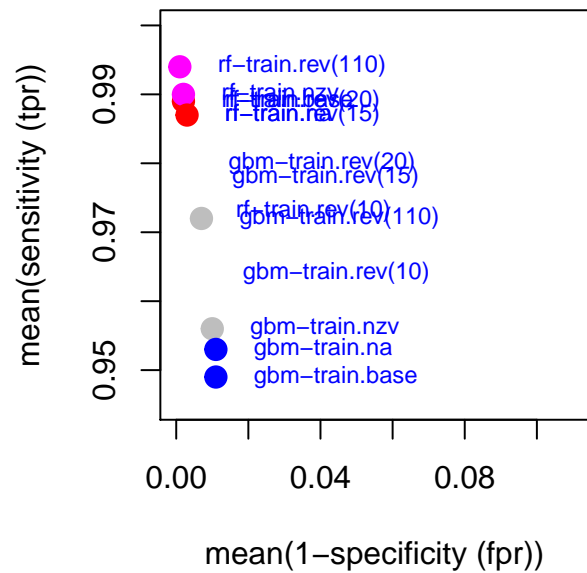
    # create output table
    res=as.data.frame(t(as.matrix(c(algos$algo[[k]],dset$ds[[i]]))))
    colnames(res)=c("algo","data set")
    sense=round(t(confusionMatrix(preds,testing$classe)$byClass[,1]),3)
    spec=round(t(confusionMatrix(preds,testing$classe)$byClass[,2]),3)
    res=cbind(res,sense,round(mean(sense),3),spec,round(mean(spec),3))
    if(exists("res.mstr")){res.mstr=rbind(res.mstr,res)}else{res.mstr=res}
    print(res.mstr)
  }
}
stopCluster(clust)

```

**'ROC' – all algos**



**'ROC' – rf/gbm**



#### top 20 vars

Finally, I re-ran the rf and gbm algorithms using the train.rev database and just their top variables. I compared four sets of results for different numbers of top variables including 100, 20, 15 and 10. Interestingly the performance did not degrade significantly for either algorithm. While rf moved from 99.5% to 97.5% accuracy between the 100 variable and 10 variable models, gbm actually increased its accuracy from 97% to 98% using as variables decreased from 100 to 15.

**'ROC' – rf/gbm**

